

Near Complete Formal Semantics of X86-64

Abstract

To Do

1. Introduction

2. Challenges

2.1. Using Strata Results

Following are the challenges in using *Strata* [4] (or *Stoke*) formula as is.

- *Stoke* uses C++-functions which define the semantics of instructions. For example, following is the function to define the semantics of add instruction. The functions are generic in the sense that they can be used to obtain the concrete semantics of any instruction like `add %rax, %rbx`
- ```

S1. void add(SymBitVector dest, SymBitVector a,
 SymBitVector b) {
S2. set(d, a+b);
S3. }
```

The untested assumption here is the generic formula will behave identically for all the variants. We have tested all the formula for each instruction variant.

- *Strata* gives the concrete semantics for a concrete instructions. For other variants it generalize from the concrete semantics. Assumption is the generalization is correct. Test all the generalization.
- While porting to K rule, we generalize the from a concrete semantics that *strata* provides. Is this generalization faithful? For instruction like `xchg, xadd, cmpxchg`, the formula is different for different operands. So the general K rule we obtain from `xchgl a, b` may not represent the semantics for `xchgl a, a`. Fortunately there exists different instruction variants if the their semantics might be different and accordingly we might have different K rules. For example, `xchgl_r32_eax` and `xchgl_r32_r32`. But even for `xchgl_r32_r32` semantics could be different for cases  $r1 \neq r2$  and  $r1 == r2$ . Idea: Once lifted as K rule, test the instruction for all variants.

Lets consider `xaddb SRC, DEST`, as per manual the semantics is as follows:

```

S1. Temp = Src + Dest
S2. Src = Dest
S3. Dest = Temp
```

The point to note here is that the register updates follow an order. *Strata* uses `xaddb %rax, %rbx`, to obtain the semantics and it happened that the ordering is maintained and hence *strata* can generalize the semantics of `xaddb R1, R1`. But even if the ordering is not maintained the semantics is going to

be the same for the case  $R1 \neq R2$ , but the generalization for the  $R1 == R1$  case will mess up. We cannot trust the above generalization by *strata*. We need to test the K rule for all possible operands.

## 3. X86-64 Instruction Semantics in K

### 3.1. Modeling Instruction Semantics

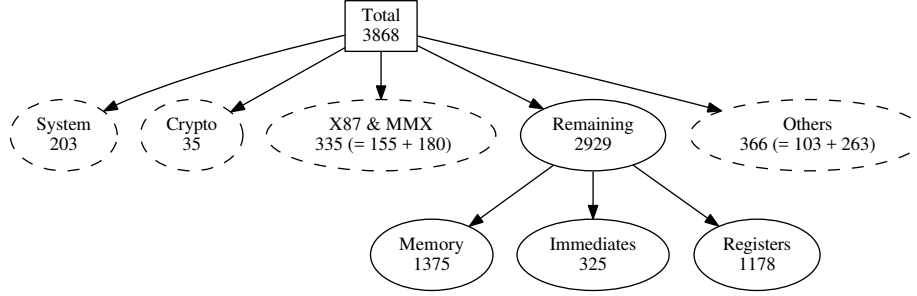
In this work we supported formal semantics of the input/output behavior of 2929 out of 3868 x86-64 Haswell ISA instruction variants. Figure 1 shows the classification of the instructions not supported using dotted ovals. **Why they are not supported??**

In order to get semantics of individual instructions, we build on top of project *Strata* [4] which automatically synthesized formal semantics of the input/output behavior for 1796 Haswell ISA X86-64 instructions. The key to their results is stratified synthesis, where they use a set of instructions whose semantics are known to synthesize the semantics of additional instructions whose semantics are unknown. Using a combination of stochastic search + pruning using testing (we refer as *initial search*) and subsequent refining of the search results using equivalence checking (referred as *secondary searches*), they first came up with the semantics of 692 register and  $\sim 120$  immediate instructions. The rest  $\sim 984$  are the immediate and memory variants obtained by generalization of 692 register instructions.

#### Strata Vs Stoke

Following are some of the immediate challenges that we needed to address.

1. **CH.1: Supporting un-stratified Instructions** The paper [4] mentions that adding some primitive instructions (like saturated add) as the base instruction might help stratified more instructions. We would like to explore similar directions. Moreover, it would interesting if we can leverage the manually written instruction semantics from project *Stoke*.
2. **CH.2: Getting Generic Formula for immediates** The  $\sim 120$  immediate instructions, mentioned above, do not have a corresponding register-only instruction to generalized from. Therefore *Strata* tries to learn a separate formula for every possible value of the 8 bit immediate operand. We intend to have a more intuitive generic semantics (that works for all values of the immediate operand) for those instructions.
3. **CH.3: Modeling undef flags** There are instructions which conditionally sets some cpu flags to *undef*. For example, the shift left instruction `salq %cl, %rbx` sets flag *%of* to



**Figure 1: Instruction classification.**

The solid ovals are the ones modeled by this work.

*undef* state if the count mask  $> 1$ . Also there are instructions like `btsr %eax, %ebx` which un-conditionally puts flags like `%pf` & `%af` into *undef* state.

*Strata* while doing the *initial search* does not test the flags which *may* (for conditional *undefs*) or *must* (for un-conditional *undefs*) be taking undefined values. We intend to model the semantics of these flags with the same correctness guarantee as the other registers which do not result in *undef* and hence modeled by *Strata*.

4. **CH.4: Modeling `%af` flag** *Strata* chose not to model the `%af` flag as this is not commonly used. Supporting this flag fall within the scope of our work.
5. **CH.5: Generalization to Immediate and Memory** How reliable is the generalization of register instructions to memory or immediate variants? *Strata* states that the claim for the generalization is based on random testing.

Following is a key observation concerning stratification which help us handle the most of the above mentioned challenges.

**Observation** In order to get the semantics of a target instruction  $I$ , *Strata* uses *Stoke* along with a set  $TS$  of 6580 test cases to synthesize an instruction sequence which agrees with  $I$  on  $TS$  (which means the output behavior of the instruction sequence matches with that on real hardware for input  $TS$ ). After having that *initial search*, they keep on searching additional sequences, called *secondary searches*, each agreeing with  $I$  on  $TS$ , in a hope of getting one which would prove non-equivalent to existing ones and thereby gaining more confidence on the search and probably an augmented test-suite (as  $TS$  might get augmented with a counter example from equivalence checker in the event of non-equivalence).

One unfavorable possibility for *Strata* is when all subsequent secondary search results proves equivalent to the one obtained from initial search and hence there are no conflicts among searches, in which case it means that secondary searches fail to add any “confidence” to the initial search result and end up giving the same correctness guarantee as provided by the initial search result. Even though in such unfavorable case, the secondary searches might have provided “better” choices to pick the final formula from. A better choice of formula do not contain uninterpreted functions or non-linear

arithmetics and are simple.

In the paper[4], it is mentioned that there are only 50 cases, where they found a (valid) counterexample. That means, there are  $762 = (692 + 120 - 50)$  instructions, whose the initial stoke search using augmented test-suite, containing 6630 ( $= 6580 + 50$ ) tests, is sufficient enough to provide a semantics with the same correctness guarantee which *Strata* provides. In other words, in most of the cases, the correctness guarantee of secondary searches is same as that of the initial stoke search using the augmented test-suite (henceforth referred as *ATS*) which *Strata* ends up with.

**Handling CH.1** For an unsupported instruction  $I$ , we either model its semantics manually or borrowed it from project *Stoke*. Once we have this candidate, we test it against hardware using *ATS*. Once the test passes we claim (from the above observation) the semantics to have the same correctness guarantee which *Strata* provides for most of its cases. This helped us finding instruction semantics bugs in Intel Manual [1] and *Stoke* [3].

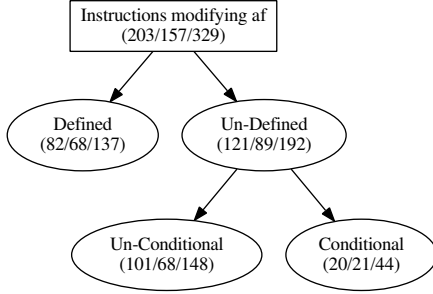
We understand that this is not as efficient as *Stoke*, which is fully automatic in getting these formulas, and we do not intend to make any contribution towards efficient generation of instruction semantics. The purpose of above mentioned effort is to deliver in cases where *Stoke* cannot without loosing much on the correctness guarantee.

Moreover, writing the semantics manually might alleviate the need of secondary search as a means to provide “better” formula as we can control the complexity and choice of operations to include in the formula. Also carefully written manual formula tend to need less number of conflicting searches than the ones generated by random search engines like *Stoke*.

differential testing add base + reduce search space

**Handling CH.2** The instructions in this category either have a separate formula for all or some of 256 possible values. We refer each of the separate formulas for instruction  $I$  as a concrete formula  $F_c^I$  for a particular constant value  $c$  of immediate operand. In either case, we get a generic formula,  $G^I$  either by writing it manually or borrowing it from *Stoke* project.

In the case where we have a separate  $F_c^I \forall c \in \{0..255\}$ , we do a Z3 equivalence check as follows:  $\forall c \in \{0..255\} : F_c^I \equiv_{Z3}$



**Figure 2: Instructions affecting %af flag.**

The numbers represent count of  
(Register/Immediate/Memory) Instructions.

$G[c]^I$ , where  $G[c]^I$  is obtained by replacing the symbolic inputs of  $G^I$  with constant value  $c$ . A successful equivalence check suggest  $G^I$  to be a generic formula with the same correctness guarantee that *Strata* has for any of the individual concrete formulas. For the case where we have a separate formula for a subset of constant values, we do the same equivalence check as before for that subset. The constants for which we do not have a separate formula we test  $G[c]^I$  using *TS*, the final test-suite of *Strata*.

**Handling CH.3** There are 474 (= 141(Reg) + 109(Imm) + 224(Mem)) instructions that results in conditional (or *may undef* (162 (= 40 + 46 + 76)) or un-conditional (or *must undef* (312 (= 101 + 63 + 148))) in one or more cpu flag. The semantics of most of the cpu flags (which *may* or *must* take *undef* values) are already modeled in *Stoke*. We needed to model the semantics of flag registers for 40 instructions involving shifts, rotates [2].

For *may undef* cases, we tested against hardware, using *TS*, for the scenarios when the condition for undefinedness is not triggered. For the remaining cases, (1) *may undefs* where the condition for triggering *undef* is true and (2) *must undef*, we make sure that *K* execution halts when the undefinedness condition is triggered. This help is find bugs in the *Stoke* implementation of 8 instructions [2] ( Note that these 8 instructions are not stratified and hence we borrowed it from *Stoke*).

**Handling CH.4** Figure 2 represents the distribution of instructions affecting the %af flag in a defined or un-defined way (which could be conditional or un-conditional). We tested all the instructions for the defined cases using *TS*. For conditionally undefined cases, we tested for the scenarios when the condition for undefinedness is not triggered. For all remaining cases, we make sure that the *K* execution halts when the undefinedness condition is triggered.

### 3.2. Porting to K Rules

For the purpose of getting *K* rules, we could have directly converted the *Strata* formulas for an instruction to *K* rule assuming that the *Strata*'s symbolic execution over the stratified instruction sequence is correct.

Given that fact the *K*'s symbolic execution engine is more trusted as that has been used extensively in language-agnostic manner to perform symbolic execution, we decided to use *K*'s symbolic executor. Also in order to check if *Strata*'s symbolic execution engine is correct, we did an equivalence check on the outputs of both the symbolic executions.

1. Implementing the base instructions semantics in *K* and testing them.
2. Symbolic execution of the stratified instruction sequences.
3. Dealing with scratch pad registers.
4. Equivalence check between *Strata* formula and the output of 2.

All the checks are *unsat*, except one where the check fail to due a bug in the simplification rules in *Strata*, which states the following lemma related to two single precision floating point numbers *A* and *B*, which is not correct for NaNs. However this bug is fixed in the latest version of *Stoke*.

$\text{sub\_single}(A, B) \equiv 0 \text{ if } A == B$

5. Simplification of formulas: Simplification generates simple *K* rule (sometimes simpler than the corresponding *Strata* formula). Also it is much easier to write the simplification rules in *K*. show the example for  $\text{concat}(A[1:2], \text{concat}(B[2:3], X)) \equiv \text{concat}(A[1:3], X)$
6. One drawback of the *Strata* formulas is they could be non-intuitive and complex at times when the simplification rules are not adequate enough to simplify their complexity to more intuitive formulas. Appendix A provides such an example. Towards the goal of having intuitive formulas, we borrowed the hand written formula (provided they are simpler) from *Stoke* or manually write those and check equivalence with the stratified formula. If they match on all register state and/or memory, we employ that in our *K* semantics.

### 3.3. Supporting un-stratified instructions & Porting their formulas to K Rules

#### 3.3.1. Supporting un-stratified instructions

##### Instruction support status

#### 3.4. Porting to K Rules

*Strata* could output the internal AST, used to model a register state formula, in different formats. Supported backend are *SmtLib* and Prefix notation. We have added another backend to generate *K* rule. We need some way to validate the backend.

**Validate the Backend** The *K* rules generated using the backend are matched (syntactically) against the ones we already obtained via symbolic execution on stratified instructions. Other than validating the backend, this has an added benefit that in order to get the exact match, we need to port all the simplification rules from *K* to *strata* code, which in turn will later help in generating simplified *K* rules for non-stratified instructions.

Main challenges in getting an exact match are:

- Strata rules uses *extract* to extract portion of a bit-vector. The high and low indices of *extract* are obtained considering LSB at index 0, whereas K uses *extractMInt* for the same purpose, but uses MSB at index zero.
- Strata uses flags as Bool, whereas they are treated as Bitvector in our semantics. We modified strata so as to treat flag registers as 1 bit bitvectors.

## A. An Example of Strata Formula

Following is the Strata formula for an instruction `vpxor %ymm3, %ymm2, %ymm1,`

```

1 (let ((a!1 (bv xor ((_ extract 255 192) ymm3)
2 ((_ extract 255 192) ymm2)
3 (bvor ((_ extract 255 192) ymm3)
4 (bv xor ((_ extract 255
5 192) ymm3)
6 ((_ extract 255
7 192) ymm2))))))
8 (a!3 (bv xor ((_ extract 191 128) ymm3)
9 ((_ extract 191 128) ymm2)
10 (bvor ((_ extract 191 128) ymm3)
11 (bv xor ((_ extract 191
12 128) ymm3)
13 ((_ extract 191
14 128) ymm2))))))
15 (a!5 (bv xor ((_ extract 127 64) ymm3)
16 ((_ extract 127 64) ymm2)
17 (bvor ((_ extract 127 64) ymm3)
18 (bv xor ((_ extract 127 64)
19 ymm3)
20 ((_ extract 127 64)
21 ymm2))))))
22 (a!7 (bv xor ((_ extract 63 0) ymm3)
23 ((_ extract 63 0) ymm2)
24 (bvor ((_ extract 63 0) ymm3)
25 (bv xor ((_ extract 63 0)
26 ymm3)
27 ((_ extract 63 0)
28 ymm2))))))
29 (let ((a!2 (bv xor ((_ extract 255 192) ymm3)
30 ((_ extract 255 192) ymm2)
31 (bvor ((_ extract 255 192) ymm3)
32 (bv xor ((_ extract 255
33 192) ymm3)
34 ((_ extract 255
35 192) ymm2))))))
36 (bvor a!1
37 ((_ extract 255 192) ymm2)
38 ((_ extract 255 192) ymm3)
39)))
40 (a!4 (bv xor ((_ extract 191 128) ymm3)
41 ((_ extract 191 128) ymm2)
42 (bvor ((_ extract 191 128) ymm3)
43 (bv xor ((_ extract 191
44 128) ymm3)
45 ((_ extract 191
46 128) ymm2))))))
47 (bvor a!3
48 ((_ extract 191 128) ymm2)
49 ((_ extract 191 128) ymm3)
50)))
51 (a!6 (bv xor ((_ extract 127 64) ymm3)
52 ((_ extract 127 64) ymm2)
53 (bvor ((_ extract 127 64) ymm3)
54 (bv xor ((_ extract 127 64)
55 ymm3)
56 ((_ extract 127 64)
57 ymm2))))))
58 (bvor a!5
59 ((_ extract 127 64) ymm3)
60 ((_ extract 127 64) ymm2)
61)))
62 (bvor a!7
63 ((_ extract 63 0) ymm3)
64 ((_ extract 63 0) ymm2)
65)))
66 (concat a!2 a!4 a!6 a!8)))

```

```

41 ((_ extract 127 64)
42 ymm2)))
43 (bvor ((_ extract 127 64) ymm2)
44 ((_ extract 127 64) ymm3) a!5)))
45 (a!8 (bv xor ((_ extract 63 0) ymm3)
46 ((_ extract 63 0) ymm2)
47 (bvor ((_ extract 63 0) ymm3)
48 (bv xor ((_ extract 63 0)
49 ymm3)
50 ((_ extract 63 0)
51 ymm2))))))
52 (bvor ((_ extract 63 0) ymm2) ((
53 _ extract 63 0) ymm3) a!7))))))
54 (concat a!2 a!4 a!6 a!8)))

```

where as following is the formula obtained from Stoke (hand-written) and Z3 took 88.70 secs to prove that they are equivalent.

```

1 %ymm1 : (bv xor %ymm2 %ymm3)

```

## References

- [1] Bug reported in Intel Developer Zone: Possible errors in instruction semantics. <https://software.intel.com/en-us/forums/intel-isa-extensions/topic/773342>, April 2018. Last accessed.
- [2] Bug reported in Stoke: Modelling the behavior of flags which may or must take undef values. <https://github.com/StanfordPL/stoke/issues/986>, May 2018. Last accessed.
- [3] Bug reported in Stoke: Semantic bugs. <https://github.com/StanfordPL/stoke/issues/983>, April 2018. Last accessed.
- [4] Stefan Heule, Eric Schkufza, Rahul Sharma, and Alex Aiken. Stratified synthesis: Automatically learning the x86-64 instruction set. In *Proceedings of the 37th ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI '16*, pages 237–250, New York, NY, USA, 2016. ACM.