

Near Complete Formal Semantics of X86-64

Abstract

ToDo

1. Introduction

2. Challenges

2.1. Using Strata Results

Following are the challenges in using *Strata* [1] (or *Stoke*) formula as is.

- *Stoke* uses C++-functions which define the semantics of instructions. For example, following is the function to define the semantics of add instruction. The functions are generic in the sense that they can be used to obtain the concrete semantics of any instruction like `add %rax, %rbx`

```

1
2 void add(SymBitVector dest, SymBitVector a,
3         SymBitVector b) {
4     set(d, a+b);
5 }

```

The untested assumption here is the generic formula will behave identically for all the variants. We have tested all the formula for each instruction variant.

- *Strata* gives the concrete semantics for a concrete instructions. For other variants it generalize from the concrete semantics. Assumption is the generalization is correct. Test all the generalization.
- While porting to κ rule, we generalize the from a concrete semantics that *strata* provides. Is this generalization faithful? For instruction like `xchg, xadd, cmpxchg`, the formula is different for different operands. So the general κ rule we obtain from `xchgl a, b` may not represent the semantics for `xchgl a, a`. Fortunately there exists different instruction variants if the their semantics might be different and accordingly we might have different κ rules. For example, `xchgl_r32_eax` and `xchgl_r32_r32`. But even for `xchgl_r32_r32` semantics could be different for cases $r1 \neq r2$ and $r1 == r2$. Idea: Once lifted as κ rule, test the instruction for all variants.

Lets consider `xaddb SRC, DEST`, as per manual the semantics is as follows:

```

S1. Temp = Src + Dest
S2. Src = Dest
S3. Dest = Temp

```

The point to note here is that the register updates follow an order. *Strata* uses `xaddb %rax, %rbx`, to obtain the semantics and it happened that the ordering is maintained and hence

strata can generalize the semantics of `xaddb R1, R1`. But even if the ordering is not maintained the semantics is going to be the same for the case $R1 \neq R2$, but the generalization for the $R1 == R1$ case will mess up. We cannot trust the above generalization by *strata*. We need to test the κ rule for all possible operands.

3. Implementation

3.1. Porting Formulas for stratified instructions to κ Rules

For the purpose, we could have directly converted the *Strata* formulas to κ rule assuming that the *Strata*'s symbolic execution over the stratified instruction sequence is correct.

Given that fact the κ 's symbolic execution engine is more trusted as that has been used extensively in language-agnostic manner to perform symbolic execution, we decided to use ...

- Implementing the base instructions semantics in κ and testing them.
- Symbolic execution of the stratified instruction sequences.
- Dealing with scratch pad registers.
- Simplification of formulas. Simplification generates simpler κ rule (than the corresponding *Strata* formula) in some cases.
- Check the the ported κ rules are equivalent to the *strata* formulas to begin with. This helped us in finding a bug in *Strata*'s symbolic execution engine which is anyway fixed by the authors in their recent versions.

3.2. Supporting un-stratified instructions & Porting their formulas to κ Rules

3.2.1. Supporting un-stratified instructions

Instruction support status

3.3. Porting to κ Rules

Strata could output the internal AST, used to model a register state formula, in different formats. Supported backend are *SmtLib* and Prefix notation. We have added another backend to generate κ rule. We need some way to validate the backend.

Validate the Backend The κ rules generated using the backend are matched against the ones we already obtained via symbolic execution on stratified instructions. This has an added benefit that in order to get the exact match we need to port all the verification lemma from κ to *strata* code, which will later help in generating simplified κ rules for non-stratified instructions.

Two main challenges in getting an exact match are:

- *Strata* rules uses *extract* to extract portion of a bit-vector. The high and low indices of *extract* are obtained considering

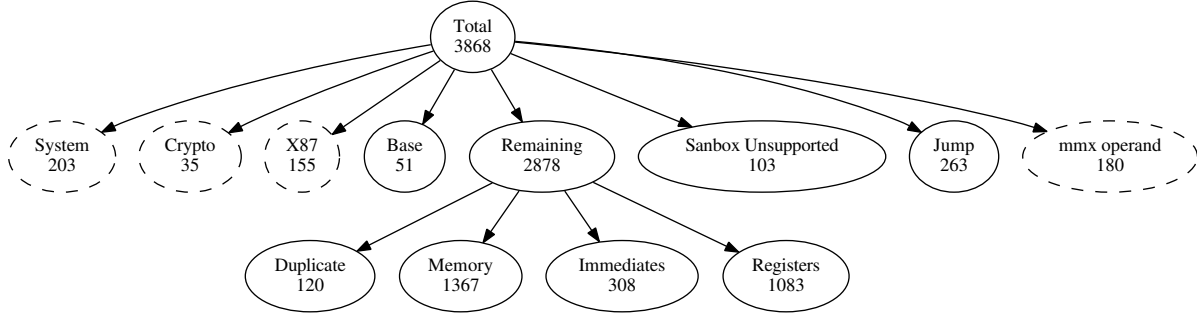


Figure 1: Instruction classification

LSB at index 0, whereas \mathbb{K} uses *extractMInt* for the same purpose, but uses MSB at index zero.

- Strata uses flags as `Bool`, whereas they are treated as `Bitvector` in our semantics.

References

- [1] Stefan Heule, Eric Schkufza, Rahul Sharma, and Alex Aiken. Stratified synthesis: Automatically learning the x86-64 instruction set. In *Proceedings of the 37th ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI '16, pages 237–250, New York, NY, USA, 2016. ACM.